

Operating Systems

LISHA/UFSC

`guto@lisha.ufsc.br`
`http://www.lisha.ufsc.br/~guto`

May 13, 2016

Computer Systems


- Hardware
 - CPU + memory + I/O devices
- Operating system
- Application programs
 - Actual goal of computer systems
 - Databases, automation, games, etc
- Users
 - Define computing problems to be solved
 - People, machines, other computers

Operating Systems

- Virtual machine perspective
 - OS extends the hardware as to implement a higher-level interface to applications
- Resource manager perspective
 - OS manages system resources (processors, memory, disk, etc) for applications' convenience

Historic Perspective

- First generation (1945 - 1955)
 - Vacuum tube
 - No software at all
 - Operated through cable switches



ENIAC (1946)




First bug 'caught' by Grace Murray Hopper, 1945.


Historic Perspective

■ Second generation (1955 - 1965)



- Transistor
- Device drivers
- First programming languages (Fortran)
- Monitor (punched card reader)



- Batch (off-line)



- Spooler (*Simultaneous Peripheral Operation On-Line*)



TX-0 Transistorized Experimental Computer (1956)

Historic Perspective

- Third generation (1965 - 1980)
 - Integrated circuit (TI IC)
 - First generic OOSS (IBM OS/360)
 - Multiprogramming (CPU/IO overlap)
 - Time-sharing (MIT CTSS)
 - MULTICS (MIT, BELL, GE)
 - PDP-11 (DEC)
 - UNIX (BELL)



PDP-11/20 (1970)

Historic Perspective

- Fourth generation (1980 - ?)
 - Microprocessor
 - MS-DOS, UNIX
 - Network systems
 - Distributed systems
 - Real-time systems




Apple MacIntosh SE/30
(1972)

Historic Perspective

- Fifth generation (?)
 - Hardware
 - Parallel?
 - Embedded?
 - Ubiquitous!
 - Software
 - Human interface!
 - Artificial intelligence???

A Typical Computer



Computer System Structures

- Motivation
 - Overlap CPU and I/O operations to improve performance
 - Avoid inter-process interference
- Interrupts
 - Avoids busy-waiting
 - I/O device receives a service request and generates an interrupt when the request has been accomplished
 - Transparent to processes
- Direct Memory Access (DMA)
 - Data transfer between I/O device and main memory without CPU assistance

Computer System Structures

- Resource protection
 - Enable the OS to define policies
 - Violations causes a *trap* into the OS
- CPU
 - Dual-mode operation
 - Supervisor mode: whole instruction set, restricted to the OS
 - User mode: unprivileged instruction subset (e.g. no I/O)
 - Timer
 - Timer interrupts periodically transfer control to OS
- Memory
 - Memory Management Unit (MMU)
 - OS isolation
 - Private address spaces for each process
 - I/O device controllers' registers protection

Operating System Services

- Process management
 - Creation and destruction of processes
 - Resource allocation and reclaiming
 - CPU scheduling (and process accounting)
 - Process synchronization
 - Process communication
 - Deadlock handling
- Memory management
 - Memory allocation and deallocation
 - Integrity maintenance (what belongs to whom)
 - Swapping
 - Virtual memory

Operating System Services

- I/O management
 - Buffering/caching
 - Scheduling (e.g. disk, network)
 - Device drivers
- File management
 - Creation, manipulation, and deletion of files
 - Creation, manipulation, and deletion of directories
 - Mapping of files onto disks
- Networking
 - Routing, contention, and security of messages
 - Heterogeneity
 - User interface

Operating System Services

- Protection
 - Access control to resources
 - Logging
 - Procedure validation
- Interface
 - OS provides services to applications by means of an Application Program Interface (API)
 - If the OS is in a different protection domain than applications (e.g. kernel), a system call mechanism is used
 - User interaction
 - Command interpreter (*shell*)

Operating System Architectures

- Monolithic
 - The whole OS comprises a single complex program that is responsible for all services
- Virtual machine
 - OS services are delivered as a private virtual machine to each application process
- Kernel + servers
 - Crucial OS parts, responsible for fundamental services, are kept in a protected kernel
 - Advanced services are delegated to servers that run as ordinary processes

Operating System Architectures

- Microkernel + servers
 - Only services needed to support server processes are kept in the kernel
- Exokernel + libraries
 - Physical resources (CPU, memory, cache) are safely exported to be handled directly by application programs
 - Libraries provide typical implementations for OS services
- Embedded into the application
 - Usually deployed with single-application systems
 - Only OS services required by the application get linked with it

Operating System Engineering

- **Structured**
 - OS is decomposed in a set of procedures
 - Modifications require the whole system to be rebuilt
- **Modular**
 - OS is decomposed in a set o modules (e.g. subsystems, service class, etc)
 - Enables replugging
- **Object-oriented**
 - Similar to modular, but using more powerful software engineering techniques
- **Component-based**
 - OS is decomposed in a set of reusable components (public interface accesses only)

Process Management

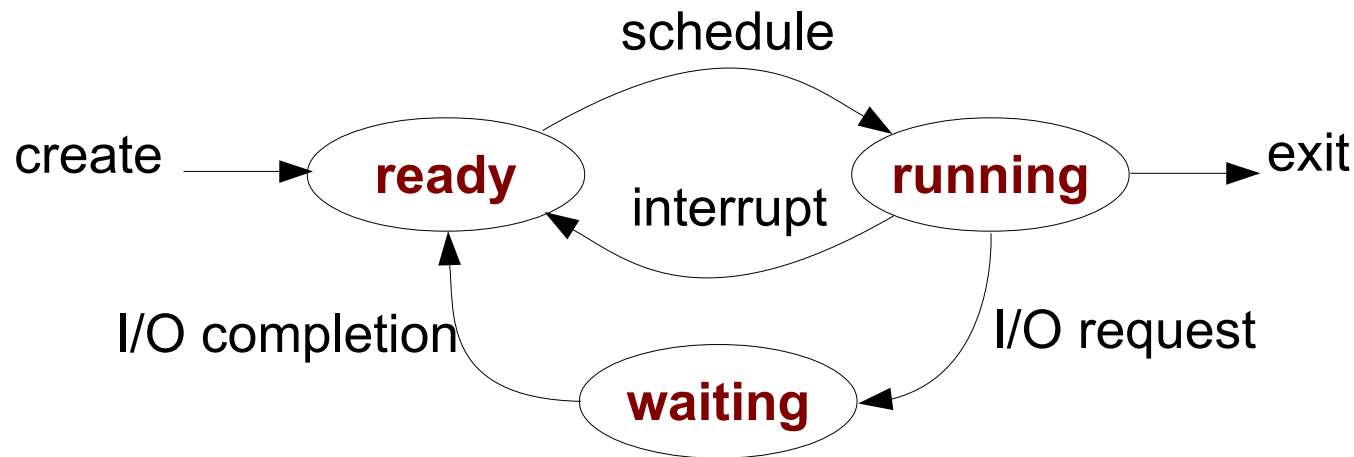
■ Process

- Is a running program
- Is an active entity
- Has context and state
- Is sequentially executed
 - A single instruction is executed on behalf of a process at any time
- Also called
 - Job on batch systems
 - Task on time-sharing systems

Process State

■ Process state


- Running: process' instructions are being executed
- Waiting: process is waiting for some event (e.g. I/O completion)
- Ready: process is ready for execution, but must waiting for a processor to become available



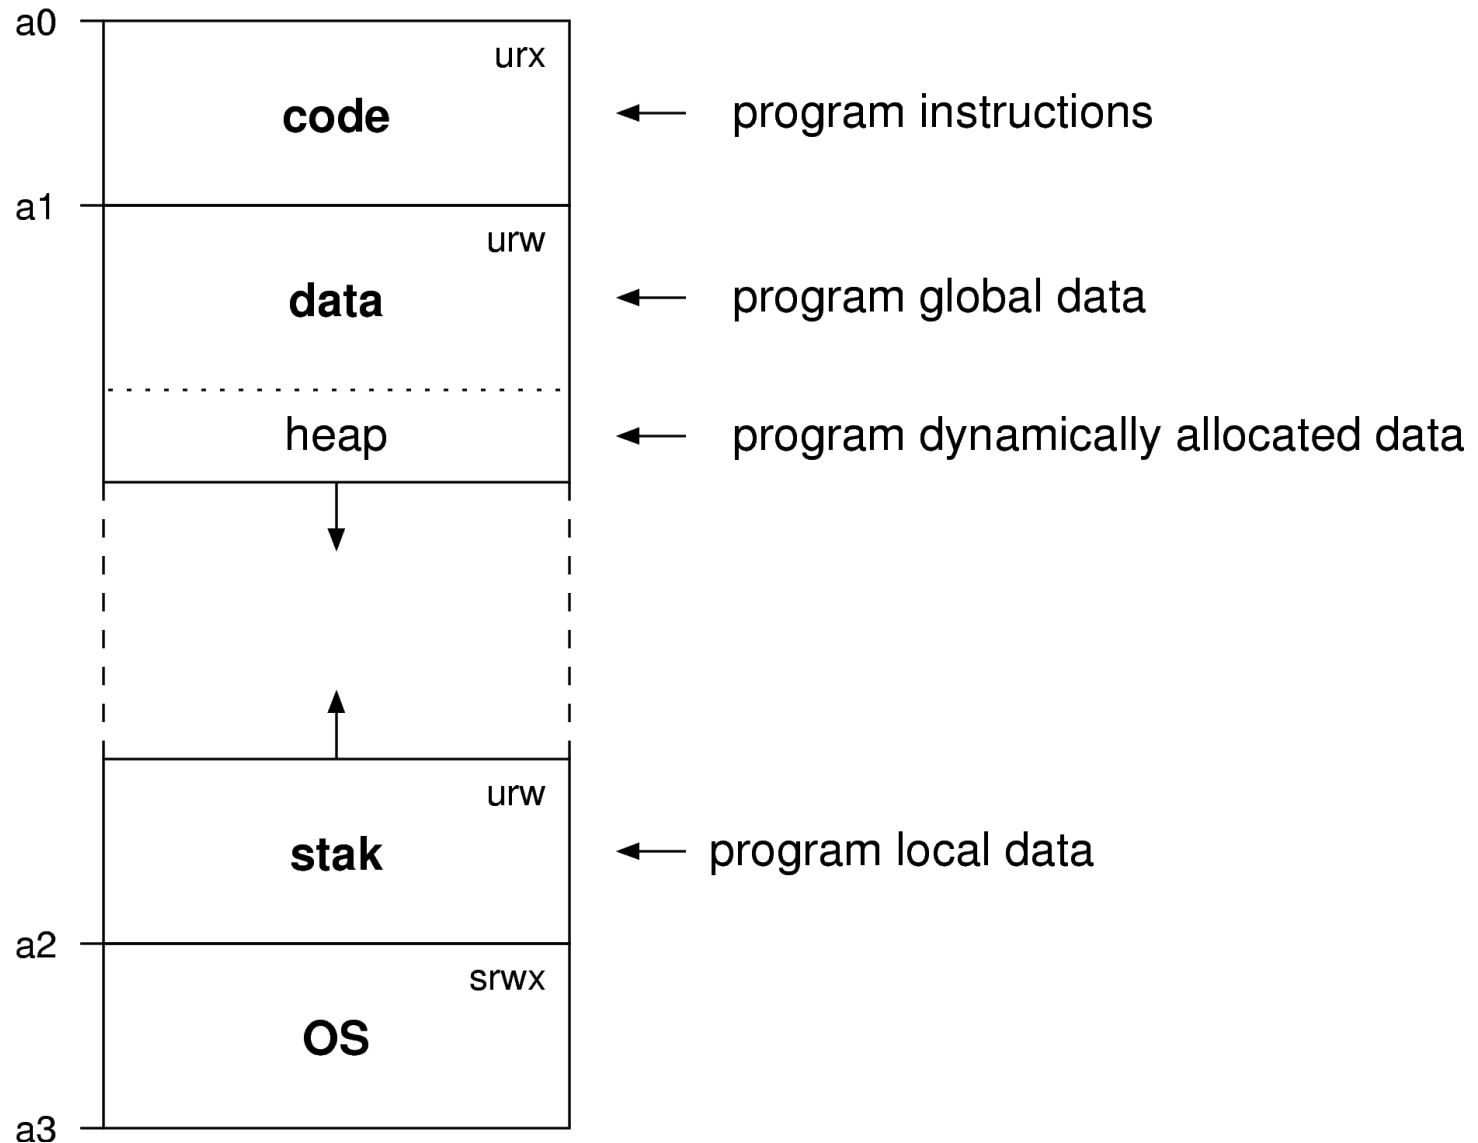
Process Context

- Process context
 - Information that allows the OS to resume the execution of a process
 - Process Control Block (PCB)
 - State
 - CPU registers
 - Scheduling info
 - Memory info
 - I/O info
 - Accounting info
 - Process stack

Context Switch



Process Address Space



Process Creation

- A process is created when another process invokes the corresponding syscall (*e.g. fork*)
 - Creator = *parent* process
 - Created = *child* process
 - Child resources can be
 - Inherited from parent
 - Allocated from OS
 - Who creates the first process?
 - Forged by OS initialization procedure
- Process destruction
 - Natural: when a process terminates and calls *exit*
 - Forced
 - By the OS when a process misbehaves (*abort*)
 - By another process (parent) on convenience (*kill*)

Concurrent Processes


- Concurrent processes
 - Resource sharing (concurrency)
 - Speedup with multiple processing elements
- Independent process
 - A sequential program under execution
 - Private context
 - Output depends exclusively from input
- Cooperating processes
 - A parallel program under execution
 - Shared context
 - Output depends also on the relative execution order

Threads

- Threads
 - Also called lightweight process
 - Low creation overhead
 - Execution flow on a task
 - Share task's code, data and resources
 - Has its own stack
 - Traditional process = task + 1 thread

Multiprogramming

■ Without multiprogramming




time = 10 tu

average execution time = 7.5 tu

throughput = 0.2 proc/tu CPU usage = 50%

■ With multiprogramming




time = 6 tu

average execution time = 5.5 tu


throughput = 0.33 proc/tu CPU usage = 100%

Process Scheduling Structures

■ Ready and I/O queues



■ System queue diagram




Process Schedulers

- Short term (CPU)
 - Selects processes from the ready queue
 - Runs very often and therefore must be very efficient
- Long term (jobs)
 - Selects processes that will be allowed in the system
 - Tries to balance I/O-bound e CPU-bound processes


Process Schedulers

- Medium-term (*swapper*)
 - Temporarily suspends processes
 - To keep the balance between I/O e CPU usage
 - Due to memory depletion



Preemptive and Non-preemptive Process Scheduling

- A process must be chosen to occupy the CPU whenever a process
 - 1 - Changes state from *running* to *waiting*
 - 2 - Changes state from *running* to *ready*
 - 3 - Changes state from *waiting* to *ready*
 - 4 - Finishes
 - Preemptive scheduling: 1, 2, 3 and 4
 - Non-preemptive scheduling: 1 and 4



Process Scheduling Criteria

- Maximize CPU utilization
- Maximize system throughput (jobs/time)
- Minimize turnaround time (total time)
- Minimize waiting time (time waiting to run)
- Minimize and stabilize (user) response time

Process Scheduling Policy

- First Come First Served
- Shortest Job First
- Static Priority
- Dynamic Priority
- Round-Robin
- Multilevel Queue
- And thousands of derivations thereof

First Come First Served (FCFS)

■ Policy


- Ready queue under FIFO policy
- New processes are inserted at the end
- Non-preemptive

■ Performance

- Extremely poor when a CPU-bound process blocks an I/O-bound process

■ Example

Process	A	B	C	D
CPU time	9	4	8	5
Arrival time	0	0	0	0



$$TA = (9 + 13 + 21 + 26) / 4 = 17.25 \text{ tu}$$

$$WT = (0 + 9 + 13 + 21) / 4 = 10.75 \text{ tu}$$


Shortest Job First (SJF)

■ Policy

- Process that will need the shortest CPU time is scheduled first
- Preemptive or non-preemptive

■ Performance

- Optimal algorithm in terms of TA and WT



$$TA = (a + (a+b) + (a+b+c) + (a+b+c+d)) / 4 = (4a + 3b + 2c + d) / 4 \text{ tu}$$


$$WT = (0 + a + (a+b) + (a+b+c)) / 4 = (3a + 2b + c) / 4 \text{ tu}$$

- Useful for processes for which the maximum execution time is known

Shortest Job First (SJF)


■ Example

Process	A	B	C	D
CPU time	9	4	8	5
Arrival time	0	1	2	3



$$TA = (9 + 12 + 24 + 15) / 4 = 15 \text{ tu}$$

$$WT = (0 + 8 + 16 + 10) / 4 = 8.5 \text{ tu}$$



$$TA = (18 + 4 + 24 + 7) / 4 = 13,25 \text{ tu}$$

$$WT = (9 + 0 + 16 + 2) / 4 = 6.75 \text{ tu}$$

SJF Approximation

■ Policy

- Future estimation based on recent past
- Process that has been having the shortest CPU cycles is scheduled first

■ Formula

$$\pi_{n+1} = \alpha t_n + (1 - \alpha) \pi_n$$

π_{n+1} = next cycle estimate

α = past importance factor


t_n = cycle n effective time

■ Example ($\alpha = 1/2$)

$$TA = (22 + 30 + 36) / 3 = 29.3 \text{ tu}$$

$$WT = (10 + 18 + 24) / 3 = 17.3 \text{ tu}$$

Process	π_0	t_0	π_1	t_1	π_2	t_2	π_3
A	1	2	1	4	2	6	4
B	1	4	2	4	3	4	3
C	1	6	3	4	3	2	2




Priority

- Policy
 - Process with highest priority is scheduled first
 - Priorities can be assigned to processes either statically or dynamically
 - Preemptive or non-preemptive
- Processes might wait indefinitely
 - Low-priority processes only run when high-priority processes are *waiting*
- Typical of real-time systems

Static Priority


■ Example

Process	A	B	C	D
CPU time	2	8	3	5
Priority	3	1	2	3
Arrival time	0	1	2	3



$$TA = (2 + 9 + 11 + 15) / 4 = 9.25 \text{ tu}$$

$$WT = (0 + 1 + 8 + 10) / 4 = 4.75 \text{ tu}$$



$$TA = (13 + 8 + 10 + 15) / 4 = 11.5 \text{ tu}$$

$$WT = (11 + 0 + 7 + 10) / 4 = 7 \text{ tu}$$


Round-Robin

- Policy
 - Processes are rescheduled periodically based on a time *quantum*
 - FIFO circular queue
 - Preemptive
- Formula
 - For a given set of processes with n elements and a time quantum of q :
 - Each process gets $1/n$ of CPU time in cycles that are no longer than q time units
 - Maximum waiting time = $(n - 1) q$
- Typical of time-sharing systems

Round-Robin

■ Example ($q = 5$ tu)

Process	A	B	C	D
CPU time	8	5	6	7
Arrival time	0	4	9	14



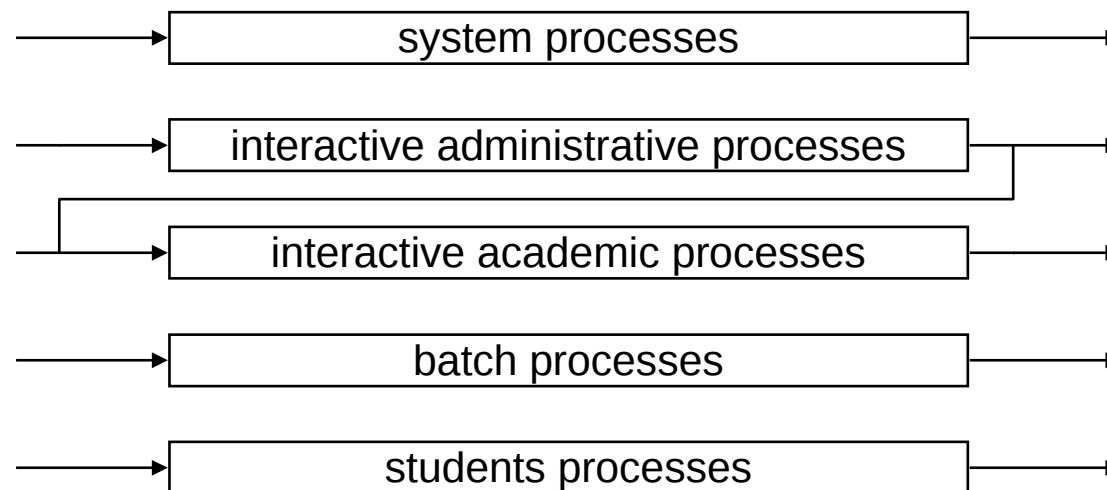
$$TA = (13 + 6 + 15 + 12) / 4 = 11.5 \text{ tu}$$

$$Wt = (5 + 1 + 9 + 5) / 4 = 5 \text{ tu}$$

Multilevel Queue

■ Policy

- Processes are grouped
 - E.g. system, interactive, batch
- Each group has its own queue under a specific policy
- Processes might be allowed to change groups



Process Synchronization

- Concurrent programs are executed by multiple cooperating processes that share some data
- Concurrent access to shared data may result in data inconsistency
- OS must provide mechanisms to synchronize and coordinate cooperating processes

Producer X Consumer

Producer:

```
shared int counter;  
shared char buf[N];  
  
int main()  
{  
    const int n = N;  
    int in = 0;  
  
    while (1) {  
        while (counter == n);  
        buf[in] = produce();  
        in = ++in % n;  
        counter++;  
    }  
}
```

Consumer:

```
shared int counter;  
shared char buf[N];  
  
int main()  
{  
    const int n = N;  
    int out = 0;  
  
    while (1) {  
        while (counter == 0);  
        consume (buf[out]);  
        out = ++out % n;  
        counter--;  
    }  
}
```

Race Conditions

Producer:

```
counter++;
load  R1,[counter]
inc    R1
store R1,[counter]
```

Consumer:

```
counter--
load  R2,[counter]
dec    R2
store R2,[counter]
```

		R1	R2	[counter]
0)	P: load R1,[counter]	5	-	5
1)	P: inc R1	6	-	5
2)	C: load R2,[counter]	6	5	5
3)	C: dec R2	6	4	5
4)	C: store R2,[counter]	6	4	4
5)	P: store R1,[counter]	6	4	6

Critical Sections

- Sections of concurrent programs in which shared data is manipulated
- Conditions for proper execution
 - Mutual Exclusion: only a single process executes a critical section on a time
 - Execution progress: a process that is not executing a critical section cannot prevent others from doing it
 - Bounded waiting: a process cannot be deprived from execution a critical section indefinitely

Synchronization Algorithm I

Process 0

```
shared int turn;

int main()
{
    while (1) {
        while(turn != 0);

        /* critical */

        turn = 1;

        /* remainder */
    }
}
```

Process 1

```
shared int turn;

int main()
{
    while (1) {
        while(turn != 1);

        /* critical */

        turn = 0;

        /* remainder */
    }
}
```

- Misses the progress condition

Synchronization Algorithm II

Process 0

```
shared int flag[2];
```

```
int main()  
{
```

```
    while (1) {  
        flag[0] = 1;  
        while(flag[1]);
```

```
        /* critical */
```

```
        flag[0] = 0;
```

```
        /* remainder */
```

```
    }  
}
```

Process 1

```
shared int flag[2];
```

```
int main()  
{
```

```
    while (1) {  
        flag[1] = 1;  
        while(flag[0]);
```

```
        /* critical */
```

```
        flag[1] = 0;
```

```
        /* remainder */
```

```
    }  
}
```

- Misses the bounded waiting condition

Synchronization Algorithm III (Peterson)

Process 0

```

shared int turn;
shared int flag[2];

int main()
{
    while (1) {
        flag[0] = 1;
        turn = 1;
        while(flag[1] &&
              turn);
        /* critical */

        flag[0] = 0;
        /* remainder */
    }
}

```

Process 1

```

shared int turn;
shared int flag[2];

int main()
{
    while (1) {
        flag[1] = 1;
        turn = 0;
        while(flag[0] &&
              !turn);
        /* critical */

        flag[1] = 0;
        /* remainder */
    }
}

```


Synchronization Hardware

■ Test and Set Lock (TSL) instruction

```
int tsl(int * ptr)
{
    int tmp = *ptr;
    *ptr = 1;
    return tmp;
}
```

■ Usage

```
shared int lock = 0;
int main()
{
    while (1) {
        while(tsl(lock));
        /* critical */
        lock = 0;
    }
}
```

Semaphores

- Integer variable accessible through atomic operations P and V

```
p(s): while(s <= 0);      v(s): s++;  
      s--;
```

- Usage

```
shared int mutex;  
int main()  
{  
    while(1) {  
        p(mutex);  
        /* critical */  
        v(mutex);  
        /* remainder */  
    }  
}
```

Semaphore Implementation

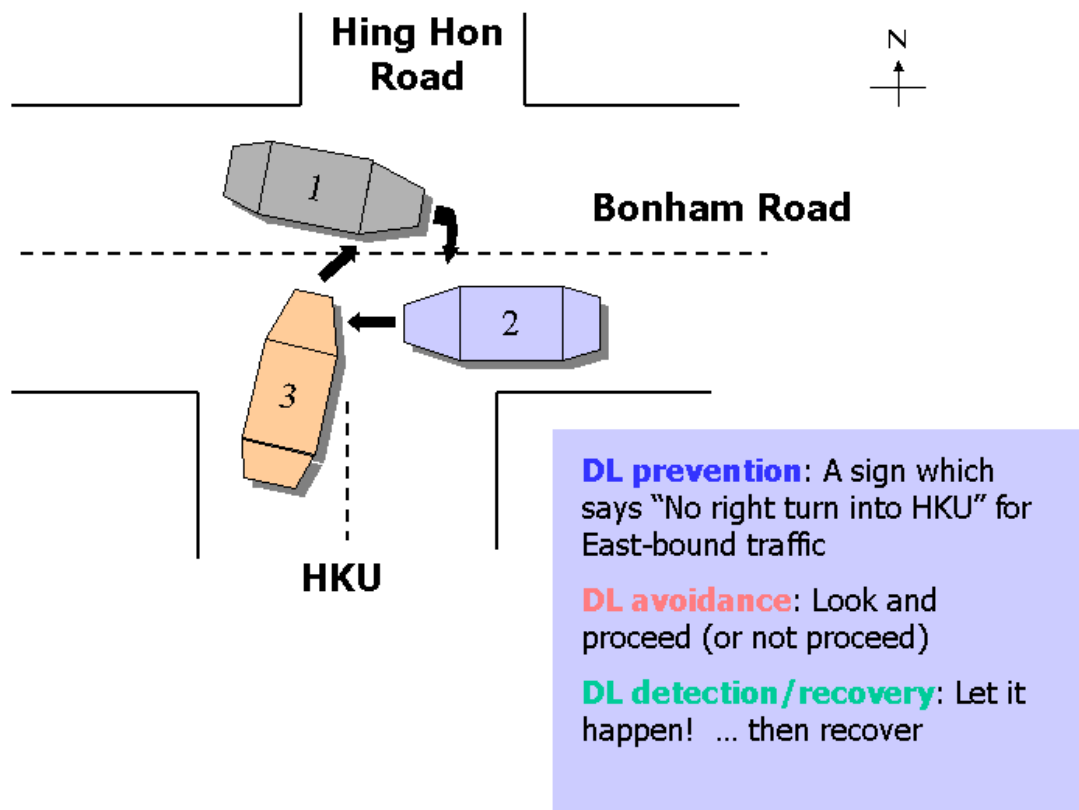
```
class Semaphore
{
public:
    Semaphore(int i) : s(i) {}
    void p();
    void v();
private:
    int s;
    list<Process> l;
};
extern Process * running;
```

```
void Semaphore::v()
{
    if(++s <= 0)
        l.pop()->wakeup();
}
```

```
void Semaphore::p()
{
    if (--s < 0) {
        l.push(running);
        running->sleep();
    }
}
```

Deadlocks

- A deadlock occurs when two or more processes are waiting for an event that can only be generated by one of the waiting processes



Deadlock Characterization

- Resource allocation
 - Request \Rightarrow Use \Rightarrow Release
- Conditions
 - Mutual exclusion: resources cannot be shared
 - Hold and wait: a process holds some resources but needs a resource that is held by another process
 - No preemption: resources cannot be preempted
 - Circular wait: there must be a circular chain of processes, each of which is waiting for a resource held by the next in the chain

Deadlock Handling


- Prevention
 - Ensure that at least one of the conditions necessary to characterize a deadlock will never hold
- Detection and recovery
 - Allows deadlocks to occur
 - Detection algorithm is run periodically
 - Allocated resources X waiting processes
 - Recovery algorithm is run whenever a deadlock is detected
 - Process termination
 - Resource preemption (rollback)
- Practice
 - Too expensive, seldom used!

Memory Management

- Processor fetches instructions from main memory
 - Programs must be loaded into memory before they can be executed
- Address referenced by programs (e.g. variables) must be bound to memory
 - Compile-time: absolute addresses
 - Load-time: relocatable code
 - Run-time: relocation hardware
- Programs bigger than memory
 - Overlays are replaced during program execution
 - Might be supported by the OS
- Often replicated programs
 - Can be organized as shared libraries

Single-Process Memory Allocation

- Without OS support
 - Simple dedicated systems (e.g. embedded)
 - No memory manager
- With OS support
 - OS memory is protected through a base register
 - User process is loaded just after OS
- Dynamic relocation with hardware support
 - Compiler and CPU issue *logic* memory addresses
 - Relocation hardware adds logic address to a base address to generate *physical* memory addresses



Multi-Process Memory Allocation without Hardware Support

- List of free memory blocks
 - First-fit: allocates the first block that is large enough to hold the process
 - Best-fit: allocates the smallest block that is large enough to hold the process
 - Worst-fit: allocates the largest available block
- External fragmentation
 - Large amount of small-size blocks
 - Enough free memory to satisfy a request but not contiguously
 - Tend to 1/3 of all the memory for n-fit algorithms

Multi-Process Memory Allocation without Hardware Support

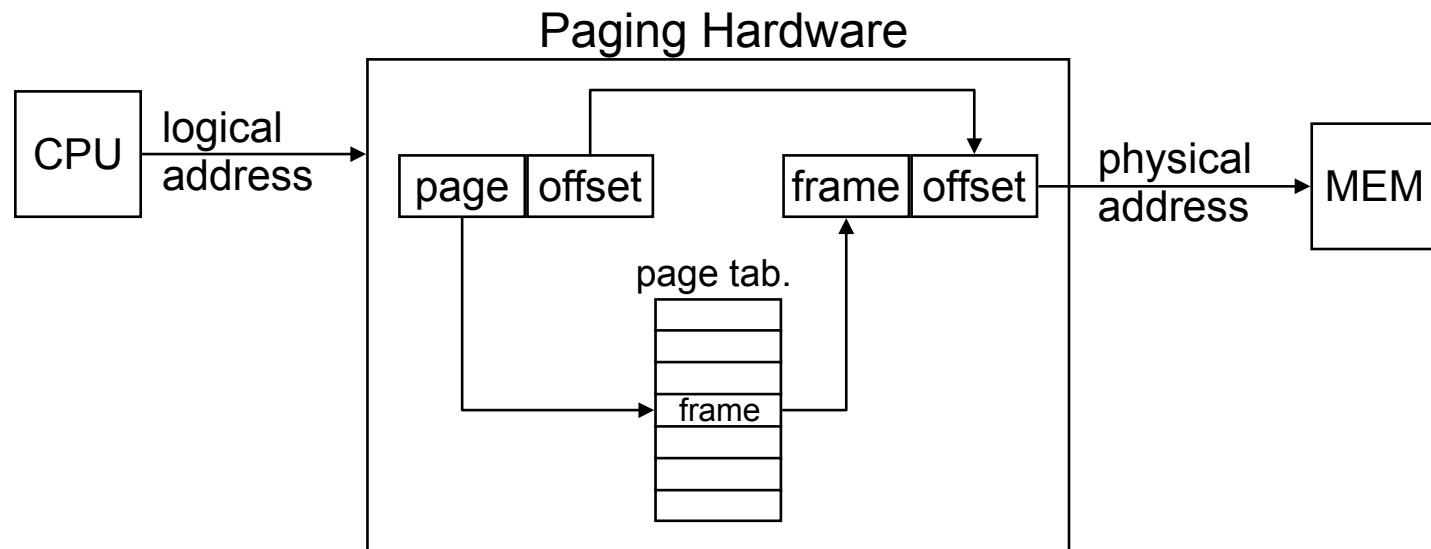
- Protection
 - Through base and limit registers
- Compaction
 - Dynamically relocates processes in order to group free blocks together
 - Relocation support
 - Relative addresses-only (implicitly relocatable)
 - Re-linking by OS application loader (expensive)
 - External relocation support

Multi-Process Memory Allocation with Hardware Support

- Detachment of address space and memory concepts
 - Compilers, processors and processes operate on an address space that is mapped into memory by an MMU
- Memory Management Unit (MMU)
 - Translates logical addresses into physical ones
 - Thus maps processes' address spaces into memory
- Typical strategies
 - Paging
 - Segmentation
 - Paged segmentation

Paging

- Memory organized and allocated in pages
- Processes address spaces
 - Organized in pages
 - Mapped into frames (physical memory pages) through page tables



Paging and Fragmentation

- No external fragmentation
- Internal fragmentation
 - Unused fraction of a page that cannot be allocated to other processes
- Internal fragmentation x page size
 - Small page size -> less fragmentation -> more memory for page tables
 - Example
 - Allocation request for 1 Gbyte
 - Pages of 4 Kbytes -> 256 Kpages
 - Pages of 4 Mbytes -> 256

Paging Implementation

- Page tables
 - Registers: limited to few pages (small address spaces or large pages)
 - Memory: slow (double memory access time)
 - Translation Look-aside Buffer (TLB)
 - Cache of page translations
 - Fast and expensive (fully associative memory)
 - Good performance if hit ratio is high (replacement policy)
- Page sharing
 - Page tables of distinct processes can reference common pages
 - Explored by shared libraries for immutable code
- Protection
 - Pages are tagged with permission bits that are checked by the MMU
 - Limit register to reduce page table size

Paging Example

Frame size: 4 Kbytes

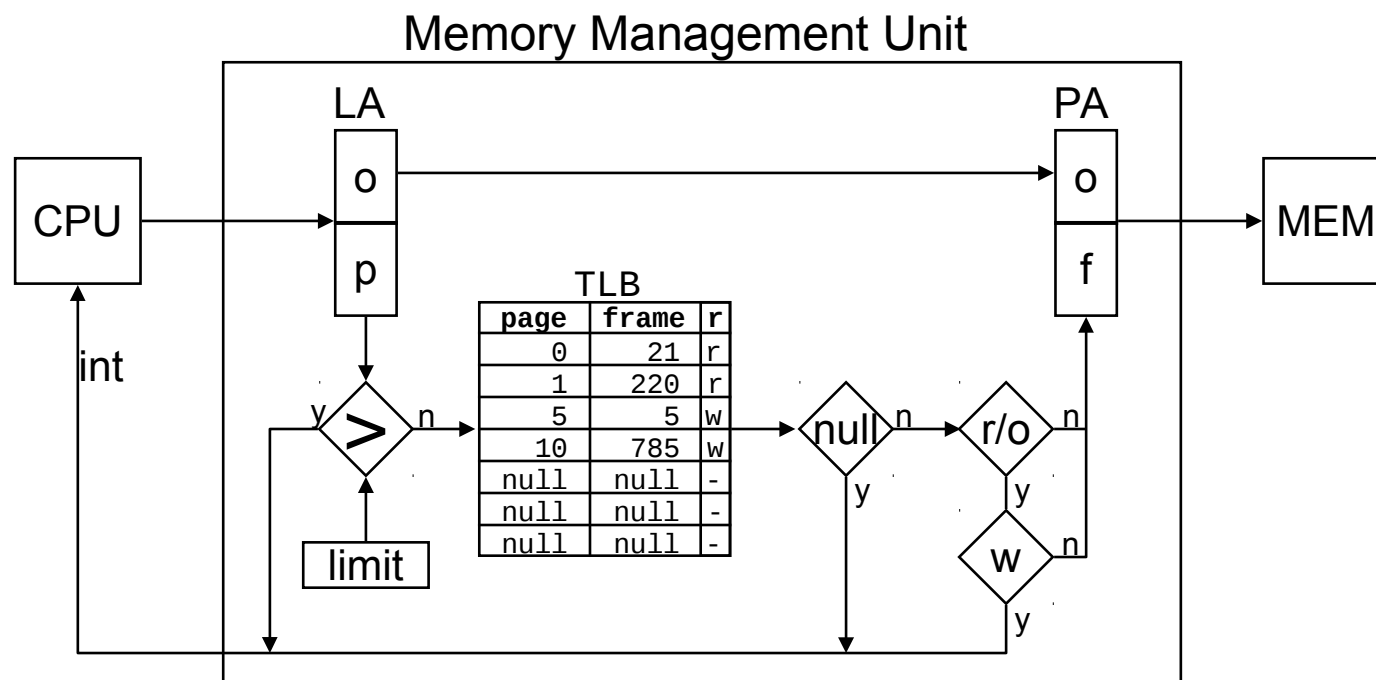
Memory size: 4 Gbytes (1 Mframes)

Address space size: 64 Mbytes (16 Kpages)

Physical address (PA): 32 bits (frame = 20, offset = 12)

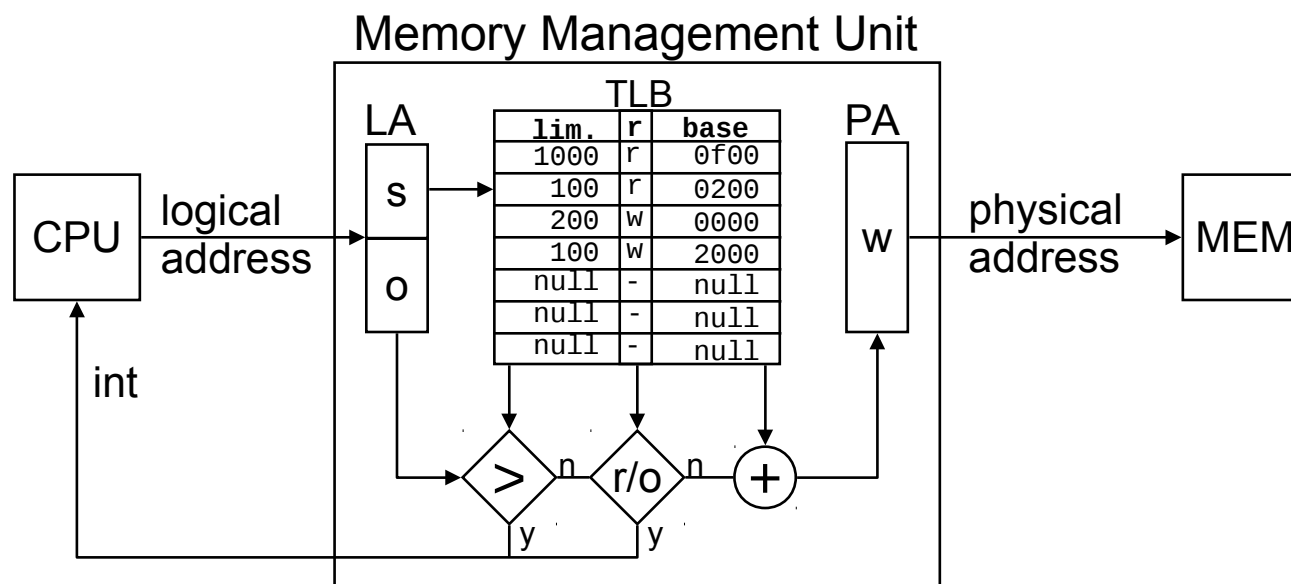
Logical address (LA): 26 bits (page = 14, offset = 12)

Page table size: 16 Kentries



Segmentation

- Memory organized in segments and allocated in words
- Bi-dimensional addresses: (segment, offset)
- Processes address spaces
 - Organized in segments
 - Mapped into physical memory through segment tables with base and limit for each segment



Segmentation and Fragmentation

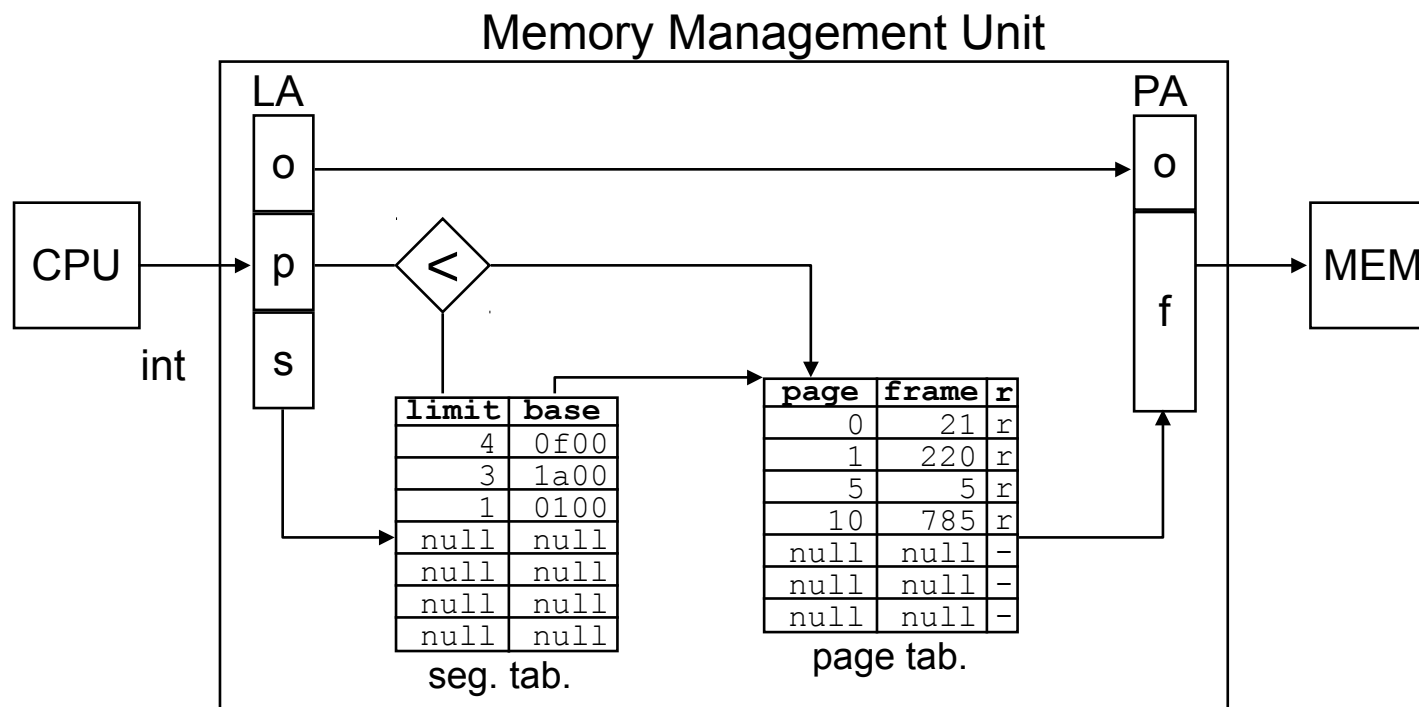
- External fragmentation
 - One segment per process -> n-fit (1/3)
 - Fixed-size segments -> paging
 - Word-size segments -> large segment tables and double memory access time
 - One segment per object
 - Intel proposal
 - Not implemented by ordinary compilers
- No internal fragmentation
 - Limit can be adjust to fit used fraction of segment

Segmentation Implementation


- Segment tables
 - Registers: limited to few segments (small address spaces or large segments)
 - Memory: slow (double memory access time)
 - Translation Look-aside Buffer (TLB)
 - Cache of address translations
 - Fast and expensive (fully associative memory)
 - Good performance if hit ratio is high (replacement policy)
- Segment sharing
 - Segment tables of distinct processes can reference common segments
- Protection
 - Segments are tagged with permission bits that are checked by the MMU
 - Segment limits are also checked by the MMU

Paged Segmentation

- Merger of segmentation and paging
- Address spaces of processes are split in segments, each of which is subsequently paged
 - Segment tables point to page tables



Virtual Memory



- Allows a process to be executed even if not completely loaded into memory
- Allows for processes to allocate more memory than the size of physical memory
- Can improve CPU utilization
- Can reduce swap overhead
- Can kill your system!

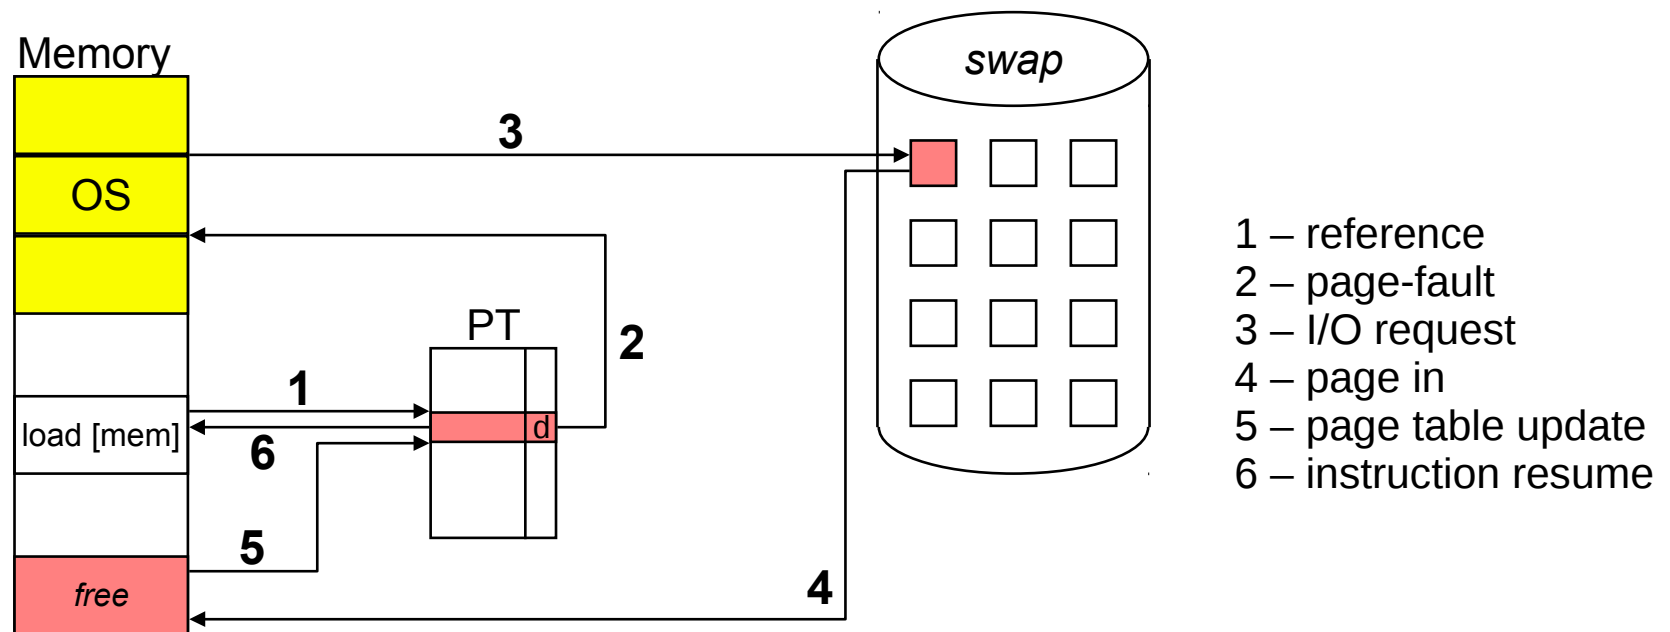
Swapping

- Processes can be temporarily suspended and the memory allocated to them is first copied to a secondary storage (i.e. disk) and then released to other processes (*swap out*)
 - Sleeping processes
 - Low-priority processes
- Such processes can be latter resumed by restoring their address spaces from the copy in the secondary storage (*swap in*)

Demand Paging

■ Page-oriented swapping

- Page table flag indicates whether the page is in memory or on disk
- MMU triggers an exception (*page-fault*) whenever an absent page is accessed



Page-Fault Handling

Page-fault trap	1 μ s
Save context	10 μ s
Dispatch PF handler	1 μ s
Locate page on disk	50 μ s
Read page from disk	10ms
Waiting queue	0s
Seek	7ms
Latency	2ms
Transfer	1ms
Scheduler	15 μ s
Disk I/O completion	1 μ s
Save context	10 μ s
Dispatch disk I/O handler	1 μ s
Update page table	15 μ s
Ready queue waiting	0s
Scheduler	15 μ s
TOTAL	10,109ms

Demand Paging Performance

■ Formula

$$\text{eat} = (1 - p) \times \text{mat} + p \times \text{pft}$$

eat = effective access time

mat = memory access time

pft = page-fault handling time

p = page-fault probability

■ Example

mat = 50 ns

pft = 10 ms = 10.000.000 ns

eat = $(1 - p) \times 50 + p \times 10.000.000$

eat = $50 + 9.999.950 \times p$

p = 0.001 \rightarrow eat = 10 μ s

eat = 50 ns $\rightarrow p < 0,000.005$ (1 / 200.000)

Page Replacement

- Whenever necessary, OS selects pages to be moved out to disk and then make them available to processes
- Algorithm criteria
 - Minimize page-fault ratio
 - Minimize I/O
- Dirty bit
 - Each page status (modified or not) is kept in the corresponding entry in the page table
 - MMU automatically sets that bit whenever a page is modified (written)
 - Modified (dirty) pages must be written to disk before being reused

First-In First-Out (FIFO)

- Replace the page that has been longer in memory (e.g. pages are time-stamped)
- Implemented using a FIFO queue
- Example

Accesses															15 pf				
7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2		2	2	4	4	4	0			0	0			7	7	7
	0	0	0		3	3	3	2	2	2			1	1			1	0	0
		1	1		1	0	0	0	3	3			3	2			2	2	1

- Increasing memory may not decrease pf rate

Accesses										9 pf		Accesses										10 pf	
1	2	3	4	1	2	5	1	2	3	4	5	1	2	3	4	1	2	5	1	2	3	4	5
1	1	1	4	4	4	5			5	5		1	1	1	1			5	5	5	5	4	4
	2	2	2	1	1	1			3	3			2	2	2			2	1	1	1	1	5
		3	3	3	2	2			2	4				3	3			3	3	2	2	2	2
														4	4			4	4	4	3	3	3

Optimal

- Replace the page that will not be used for the longest period of time
- Not implementable, for it relies on knowing the future
- Example

Accesses																	9 pf		
7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2		2		2			2			2				7		
	0	0	0		0		4			0			0				0		
		1	1		3		3			3			1				1		

Least Recently Used (LRU)

- Uses the recent past as an approximation of the near future
- Replace the page that has not been used for the longest period of time
- Example

Accesses **12 pf**

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

- Implementation
 - Time-stamp for each page
 - Linked-list of pages

LRU Approximations

- Reference bit
 - Each page is assigned a reference bit that is set by the MMU whenever the page is accessed
 - OS clears those bits periodically
 - Order of use is unknown
 - Target page is any with cleared reference bit
- Reference word
 - Additional reference bits that are shifted by the MMU
 - Target pages are those with smallest reference values
- Second chance
 - Pages are tracked by a circular FIFO list
 - If the pointed page has a clear reference bit, it is taken to be replaced
 - Otherwise, the bit is cleared and the pointer is adjusted to the next page

LRU Approximations

- Least Frequently Used (LFU)
 - Uses a reference counter for each page that is incremented by the MMU
 - Target page is the one with smallest counter value
 - Pages intensively accessed in the past, but no longer in use will take long to be replaced
- Reference and Modification bits
 - In addition to accesses, MMU marks pages that have been modified
 - Replacement order
 - Not-accessed, not-modified
 - Not-accessed, modified
 - Accessed, not-modified
 - Accessed, modified

Allocation of Frames

- Minimum set of frames
 - Instructions and operands may be scattered across several pages
 - Architecture-dependent
 - Instructions must be restarted after a page-fault
- Frames per process
 - Proportional to process size (i.e. memory footprint)
 - Equal to all processes
- Process interference
 - A process might cause the replacement of a page initially allocated to other process
 - A process may only replace its own pages

Thrashing

- A process is “thrashing” if it spends more time replacing pages than executing
- Causes
 - OS monitors CPU utilization and allows more processes in
 - If CPU utilization was low due to page-faults, increasing the number of processes in the system might cause thrashing
 - Thrashing only occurs if global page replacement (i.e. from other processes) is allowed
- Prevention
 - At any given time, a running process must have a set of pages available that fulfills its demands: its working set of pages (time x space locality)

Final Considerations

- Process load
 - On-demand
 - At-once to main memory
 - At-once to swap disk
- Page size
 - Large -> less page-fault, less I/O, less page tables
 - Small -> less internal fragmentation
- I/O results
 - Pages that will receive I/O results must be pinned-down
- Programming and code generation
 - Although virtual memory is functionally transparent to programs, memory access patterns might have big influence on it (e.g. matrices)

File Management

- Motivation
 - Common interface to transparently manipulate data on secondary storage
- File system
 - Abstract a storage device (e.g. disk) as
 - A collection of files (data) plus
 - A directory structure (control information)
 - Interaction with storage devices through services exported by the corresponding device drivers
 - Device = linear array of blocks
 - One of the most visible OS structures
 - Examples:
 - FAT, UFS, EXT2, NTFS, ISO9660, etc

Files

- File
 - Named, nonvolatile sequence of bits, bytes, lines or records
- Typed file
 - Internal structure defined by the OS
 - Executable files, graphics files, text files, etc
 - Limited number of known types
- Untyped file
 - Streams of bytes whose meaning is defined by the user
 - Unlimited and flexible

File Attributes

- Name
 - Character string identifying the file to users
- Type (only for typed files)
 - OS internal type information
- Location on device
- Size
- Ownership
- Access control
 - Who can access the file for what operations
- Access history
 - Dates, times, users, counters, etc

File Operations

- Creating
 - Locate space in the file system
 - Create a directory entry
- Deleting
 - Search the directory for the named file
 - Release file system space
 - Remove the corresponding directory entry
- Writing/reading
 - Search the directory for the named file
 - Determine the location in the file system to operate
 - Write/read data
 - Update the file pointer

File Operations

- Positioning
 - Search the directory for the named file
 - Move the file pointer
- Opening/closing
 - Since all file operations require a directory search, it is usual to implement these operations to fetch significant file's information into the system 'table of open files'
- Memory mapping
 - Associate a portion of a process' address space with a section of a file, so that reading and writing to that memory region is equivalent to performing the corresponding operations on the file

File Access Methods

- Sequential
 - Ordered access, one record after the other (tape model)
 - File pointer incremented after each operation
 - Rewind moves the file pointer back to the beginning
- Direct
 - File pointer can be moved arbitrarily (disk model)
- Indexed
 - Based on the direct access method
 - Index associating a search key to records

File Consistency Semantics

- Unix
 - Writes to an open file by a user are immediately visible to all other users that have that file open
 - Locking mechanism for access synchronization
- Session (Andrew)
 - Every new open returns a 'copy' of the file
 - No file concurrency (private copy)
 - Update at close (visible to new sessions)
- Immutable-shared-file (Bullet)
 - Shared files are made read-only

File Access Control

- Motivation
 - Multiuser file system call for access control by the OS
- Types of access
 - Read, write, execute, append, delete
- Access criteria
 - Knowing the name of files
 - Knowing a password associated to files or directories
 - Impractical for interactive applications
 - Being included in a file or directory access list
 - Associate users and access permissions
 - Hard to maintain
 - Variable size structures

File Access Control

■ Unix approach

- Simplified access list
 - Permissions to reading, writing (deleting), and executing (entering)
 - Permissions for owner, owner's group, and others
- Example
 - `drwxr-xr-x` `dir1` owner can write, all can read and navigate
 - `-rwxrwxrwx` `fil1` all can do everything
 - `-r-x-----` `fil2` owner can execute
 - `-r--r--r--` `fil3` all can read

Directories

- Directory
 - Collection of information about files
 - Translation table (name => control info)
- Device directory
 - Files' physical characteristics
 - Size, location on disk, owner, etc
- File directory
 - Volume's table of contents
 - Associate file names with device directory entries


Directory Operations

- Create/delete directories
- Add and remove directory entries
 - For file creation/deletion
- Manipulate directory entries
 - For file renaming or control information updating
- Search for a file or pattern
- Listing
- Traversing
 - For file system-wide operations such as search and backup

Directory Organization


■ Flat

- Single directory with all files



■ Tree


- OS differentiates nodes (directories) and leaves (files)
- Root node ('/')
- Pathnames
 - Absolute ('/')
 - Relative to current directory ('CWD')



Directory Organization

■ Acyclic graph


- Hard link
 - Reference counter
 - File and link are indistinguishable
 - Not applicable to directories
- Symbolic link
 - Pathname
 - File and link can be distinguished
 - May become broken
- Name aliasing problems
 - Deleting
 - Traversing



Directory Organization

■ General graph


- Cycles are allowed to exist in the directory
 - Hard links to directories
- Search algorithm must detect cycles
 - Avoid infinite loops
- Garbage collection (self reference)



Block Allocation Methods

■ Contiguous allocation

- Directory = (name, start, length)
- Optimal sequential access plus direct access
- File size defined at creation-time
 - A sufficiently large set of contiguous blocks must be located (first/best/worst-fit)
- External fragmentation
 - Garbage collection




Directory

file	start	length
f1	0	2
f2	3	4
f3	9	7

Block Allocation Methods

■ Linked allocation

- Directory = (name, start, end)
- Files are linked lists of blocks
 - Any block can be linked to any file
 - No external fragmentation
- No direct access
- Limited reliability



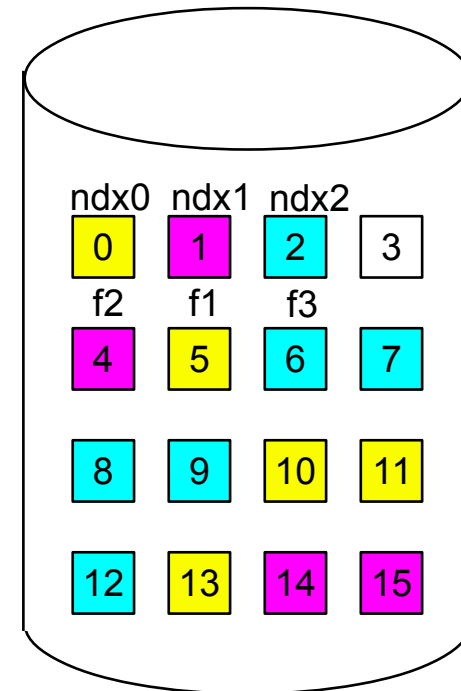
Directory

file	start	end
f1	0	7
f2	4	11
f3	15	1

Block Allocation Methods

■ Indexed allocation

- Directory = (name, index)
+ index
- Similar to paging
- Direct access without external fragmentation
- Large files
 - Linked index
 - Multilevel index



Directory

file	index
f1	0
f2	1
f3	2

index0

5
10
13
11
-

index1

4
14
15
-
-

index2

6
7
8
9
12

Free-Block Management

■ Bit map

- Each block is represented by one bit (free/used)
- Easy to locate sequences of same-state bits
 - Supports contiguous allocations
 - Optimizes sequential access
- Must reside in memory to be efficient

■ Linked list

- Free blocks are linked in a list
- Allocation and releasing imply in I/O


■ Grouping

- First free block groups a set of free blocks and contains a pointer to the next grouping block
- Contiguous ranges of blocks can be represented as pointer + count

Case Study: MS-DOS File Allocation Table (FAT)

■ MS-DOS FAT

- Directory = (name, start, end) + FAT
- Table with one entry per block is kept separately
- Special values for free blocks and end of files
- Allows direct access
- Reliability improved by replication



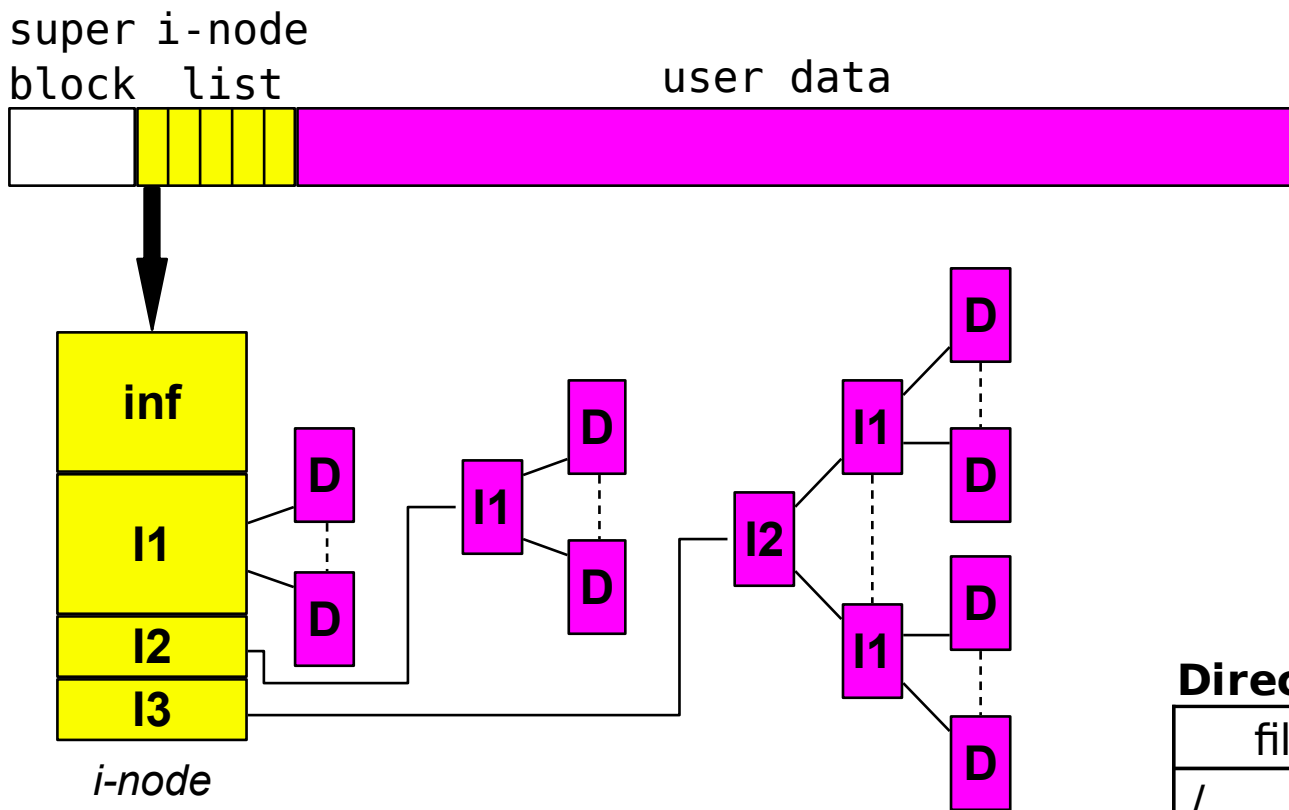
Directory

file	start	end
f1	0	7
f2	4	11
f3	15	1

FAT

2	eof	3	7
9	6	10	eof
free	5	11	eof
free	free	free	1

Case Study: Unix File System



Directory



file	i-node
/	0
f1	1
f2	2

Secondary Memory Management

- Motivation
 - Main memory is small (expensive) and volatile
 - Secondary memory is large (cheap) and persistent
 - Typically disks
- Operational basis for important OS components
 - Swapping
 - Virtual memory
 - File system


Disk Drives

- Physical structure
 - Media
 - hard or flexible, fixed or removable
 - Driver (mechanical)
 - disk rotation and head positioning
 - Controller (electronic)
 - operation and host interfacing
- Technologies
 - Magnetic
 - Optic
 - Optomagnetic



Disks

- Physical structure
 - Concentric tracks divided in sectors
 - Inter-sector gaps
 - Sectors are typically formatted to be 512 bytes-long
- Logical structure
 - Unidimensional, linear array of blocks
 - Block = 1 or n sectors
- Translation
$$\text{blk} = \text{sec} + \# \text{sec/tra} \times (\text{sur} + \text{cyl} \times \# \text{tra/cyl})$$



- Partition
 - Set of contiguous disk cylinders considered by the OS as an autonomous **logical disk**


Disk Scheduling

- Disk access time parameters
 - Seek: time to move the arm to a given cylinder
 - Latency: delay until a sector passes under the head
 - Transfer: time to transfer data from the disk controller to main memory
- Disk access requests
 - Disk address + memory address + size
 - Request queue
 - Order requests gathering those for the same cylinder
 - Order requests for different cylinders to reduce seek time
- Other performance factors
 - File organization (contiguous/disperse)
 - Control info location
 - Cache

Disk Scheduling Algorithms


■ First-Come First-Served (FCFS)

Queue: 98, 183, 37, 122, 14, 124, 65, 67 Seek: 640 tracks



■ Shortest Seek Time First (SSTF)


Queue: 98, 183, 37, 122, 14, 124, 65, 67 Seek: 236 tracks



Disk Scheduling Algorithms


■ Scan (Elevator)

Queue: 98, 183, 37, 122, 14, 124, 65, 67 Seek: 208 tracks




■ Circular Scan (C-SCAN)

Queue: 98, 183, 37, 122, 14, 124, 65, 67 Seek: 326 tracks



Redundant Array of Independent Disks

- RAID 0 (stripping)
 - Each block is broken down in sub-blocks
 - Each sub-block is stored on a different disk
 - High performance
- RAID 1 (shadowing/mirroring)
 - Each block is stored twice
 - High reliability
- RAID 5 (stripping + rotating parity)
 - High performance with good reliability



I/O Management

- Interactive systems are often more concerned with I/O than computing
- I/O devices
 - Vary widely in functionality and speed
 - Standard software and hardware interfaces help to incorporate new devices
 - New devices are constantly introduced
- Device driver
 - Bridge between OS subsystems and I/O devices
 - Encapsulate device particularities delivering an uniform interface

I/O Hardware

- **Port**
 - Host connection point for I/O devices
- **Bus**
 - Shared set of wires and a protocol that allows several devices to be simultaneously connected to the host
- **Controller**
 - Controls the operation of ports, buses and devices
 - From simple electronics to complex processors
 - Interacts with host through registers
 - Control, status, data in/out
 - I/O ports, memory mapped, CPU register mapped

I/O Operation

■ Polling

- Host 'polls' status registers to determine the status of a device
- Busy-waiting
 - Loop reading a status register
 - Overhead on multitask systems
 - Simplicity and efficiency on single-task systems

■ Interrupts

- Avoids busy-waiting
- I/O device receives a service request and generates and interrupt when the request has been accomplished
- Transparent to processes

I/O Data Transfers

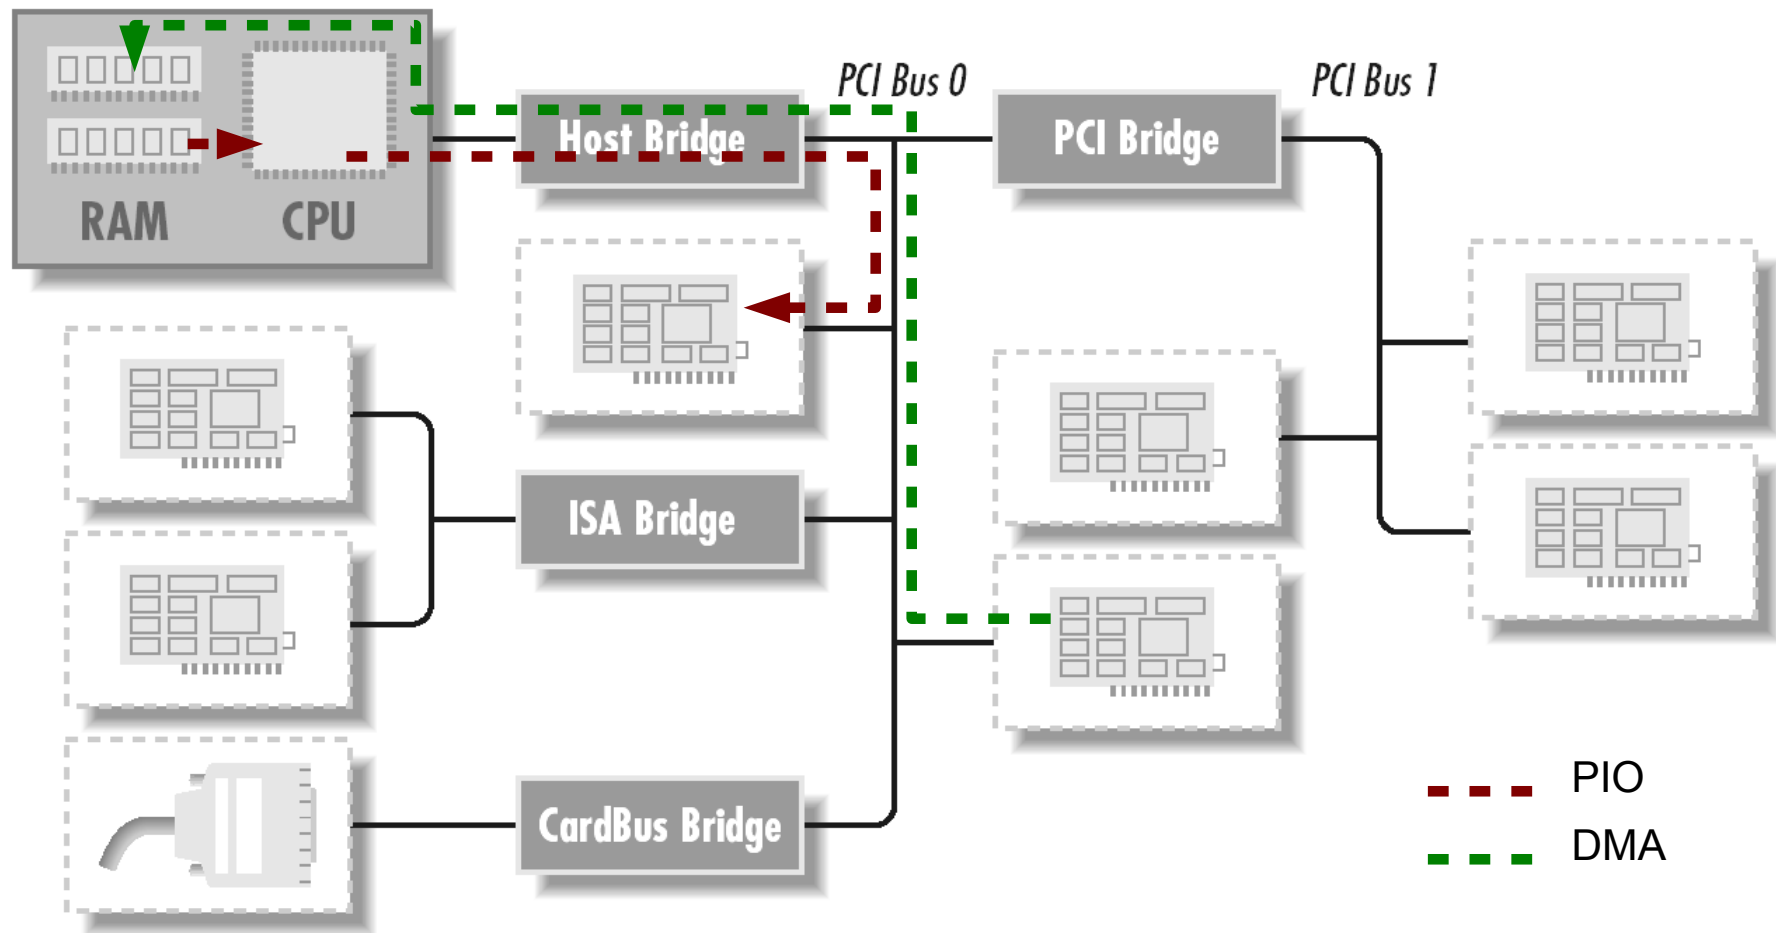
■ Programmed I/O

- Data is transferred to/from I/O device by having the CPU to write/read data registers on the device controller
- One word at a time

■ Direct Memory Access (DMA)

- Data is transferred by dedicated circuitry (DMA controller) without CPU assistance
 - Source and destination pointers + count
 - Multi-word (burst) transfers
 - Interrupt on completion or error
- Concurs with CPU for memory
- Pitfall
 - Address translation logical -> physical or DVMA

I/O Hardware



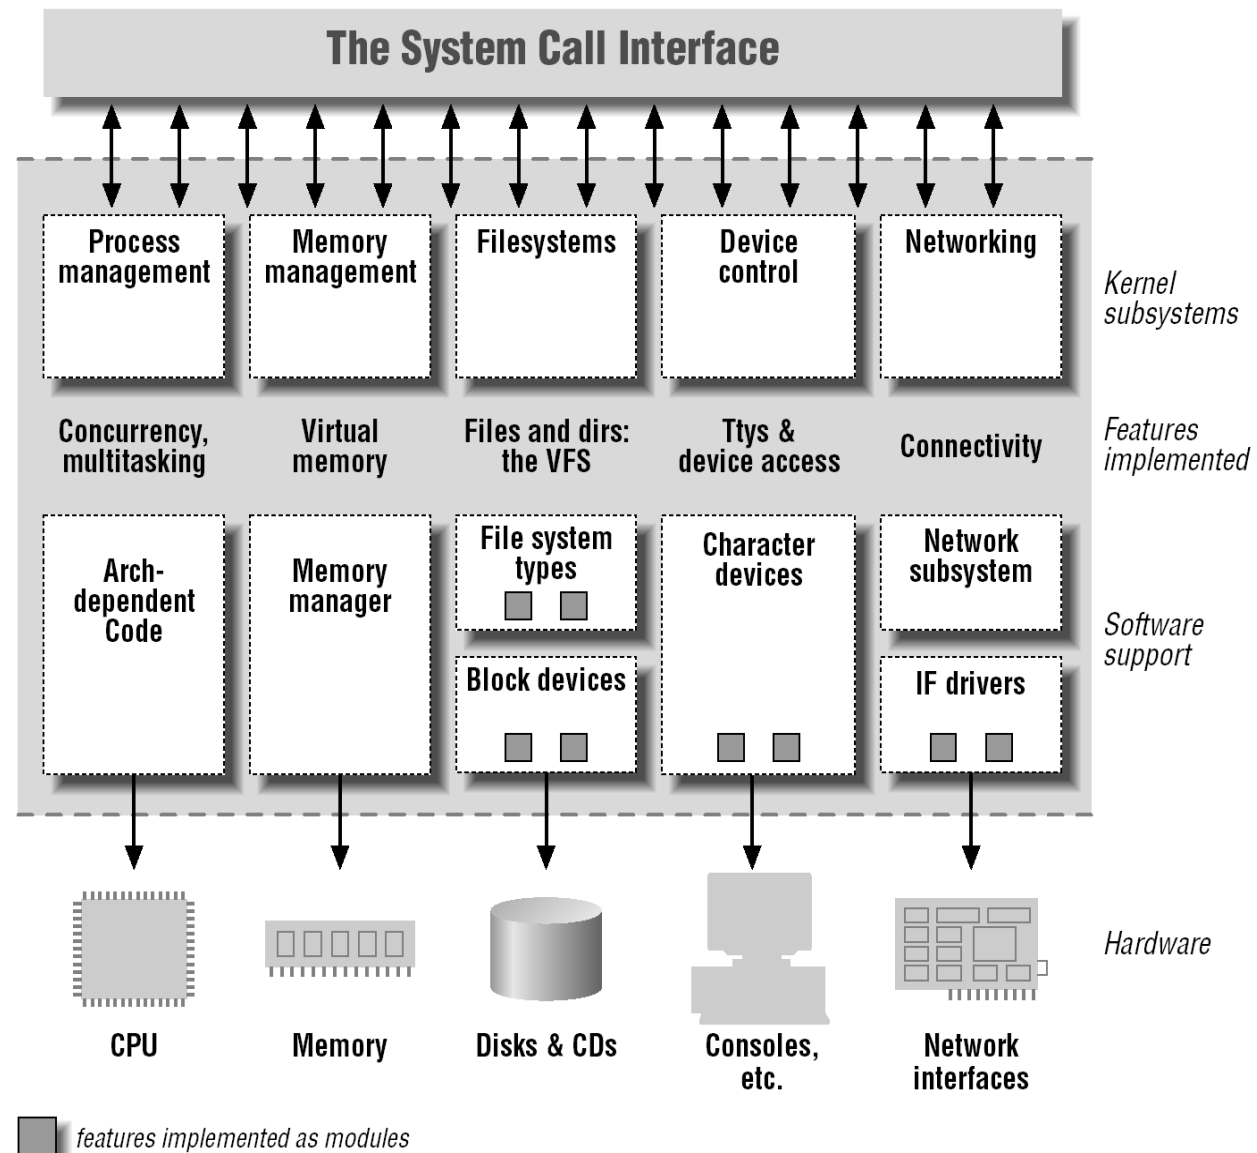
Application I/O Interface

- Indirect via I/O subsystems
 - A disk can be indirectly accessed through the files contained on it
 - A network adapter can be indirectly accessed through the TCP/IP stack (socket)
- Pseudo-file
 - Device drivers become handlers of operations on 'special files' that are plugged into the file system (/dev/mouse, /dev/hda, etc)
- Specific system calls
 - OS provides specific system calls to interact with I/O devices (eg ioctl on Unix)

Unix (Linux) Device Drivers

- Kernel module that handles the interaction with an specific hardware device, hiding its operational details behind a common interface
- Three basic categories
 - Character
 - Block
 - Network

Kernel Overview: LINUX



Hardware Devices

- Accessible via `/dev` pseudo-files
- Kernel redirect pseudo-file operations to proper device driver services considering *major* and *minor* numbers
 - *Major*
 - Identifies a driver within the kernel (8 bits)
 - *Minor*
 - Identifies a device (unit) within the driver

Char Devices

- Byte streams (e.g. `/dev/console`, `/dev/ttyS0`, `/dev/st0`)
- Operate mostly like ordinary files
 - No backward seeks

Block Devices

- Block-accessible devices at I/O level
- File system related devices (e.g. disks)
- Share a common interface with char devices, but distinct semantics
 - Block oriented (accessing single bytes is a waste)
 - Seekable
- Additional operations to support file systems

Net Devices

- Do not fit properly under the pseudo-file interface
 - Usually not a node in a file system
 - Integration with a protocol stack
- Generic network interface instead
 - Communication related operations (e.g. sending, receiving, package marshaling, time-out handling, statistic collection)
 - Optimized for TCP/IP integration

Hello World Module

```
[root]#cat > hello.c
#define MODULE
#include <linux/module.h>
int init_module(){printk("Hello World!"); return 0;}
void cleanup_module(){printk("Good Bye!");}
^D

[root]# gcc -c hello.c
[root]# insmod hello.o
[root]# dmesg
[root]# rmmod hello
[root]# dmesg
```

Module Initialization

■ Initialization

```
int init_module(void)
```

- Module's entry point
- Called at loading (by insmod)
- Performs module registration

■ Finalization

```
void cleanup_module(void)
```

- Module's exit point
- Called at unloading (by rmmod)
- Performs module unregistration

Module Registration

- Binds a module to the kernel's `syscall` interface
- Registration

```
int register_chrdev(unsigned int major, const char *name, struct file_operations *fops)
```
- Unregistration

```
int unregister_chrdev(unsigned int major, const char *name)
```
- Pseudo file

```
mknod /dev/devname0 c major minor
```

Module Parameters

- Externally accessible module-global variables
- Declared via MODULE_PARM macro

```
int irq = 10;  
char * name = "Unknown";  
MODULE_PARM(irq,"i"); /* declare irq as int */  
MODULE_PARM(name,"s"); /* declare name as string  
*/
```

- Defined at load time

```
insmod mod.o irq=9 name="The Server"
```

Module Info

- Externally visible module declarations used to supply clients with some useful information
- Macros

```
MODULE_AUTHOR("Somebody");  
MODULE_DESCRIPTION("This module doesn't do  
    anything");  
MODULE_PARM_DESC(irq, "Device IRQ (3/4)")
```

struct file_operations

```
struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char *, size_t, ...
    ssize_t (*write) (struct file *, const char *, ...
    int (*readdir) (struct file *, void *, filldir_t);
    unsigned int (*poll) (struct file *, struct ...
    int (*ioctl) (struct inode *, struct file *, ...
    int (*mmap) (struct file *, struct vm_area_struct *);
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *);
    int (*release) (struct inode *, struct file *);
    int (*fsync) (struct file *, struct dentry *, ...
    int (*fasync) (int, struct file *, int);
    int (*lock) (struct file *, int, struct file_lock *);
    :
    :
};
```

struct file

```
struct file {
    struct list_head      f_list;
    struct dentry         *f_dentry;
    struct vfsmount       *f_vfsmnt;
    struct file_operations *f_op;
    atomic_t              f_count;
    unsigned int          f_flags;
    mode_t                f_mode;
    loff_t                f_pos;
    unsigned long          f_reada, f_ramax, f_raend,
                          f_ralen, f_rawin;

    struct fown_struct    f_owner;
    unsigned int          f_uid, f_gid;
    int                   f_error;
    :
};
```

```
MINOR(f_dentry->d_inode->i_rdev)
```

Module's Reference Counter

- Automatically tracks how many clients a module has at a moment
- Avoids unloading a module that is being used by a client
- Manipulated by macros
 - MOD_INC_USE_COUNT
 - MOD_DEC_USE_COUNT
 - MOD_IN_USE

Programming Hits

- No standard libraries (and headers)
 - `printf` instead of `printk`
 - `#include <linux/x.h>`
 - `#include <asm/y.h>`
- Signalize kernel code
 - `#define __KERNEL__`
- Avoid name-clashes
 - Local symbols (`static`)
 - Prefixed symbols (`mod_sym`)
 - `EXPORT_NO_SYMBOLS;`

More Programming Hits

- Kernel code runs within the context of calling user process

```
#include <asm/uaccess.h>
```

```
unsigned long copy_to_user(to, from, count);
```

```
unsigned long copy_from_user(to, from, count);
```

- In-kernel memory allocation

```
#include <linux/malloc.h>
```

```
void *kmalloc(unsigned int size, int priority);
```

```
void kfree(void *obj);
```

Distributed Systems Taxonomy

- Stand-alone computing systems
 - Independent computers
 - Independent tasks
- Networked computing systems
 - Interconnected independent computers
 - Processes of independent tasks can communicate
- Distributed computing systems
 - Loosely-coupled computers
 - Processes of individual tasks transparently share resources
- Parallel computing systems
 - Tightly-coupled processing units
 - Several processes cooperate on a single task

A New Perspective

- **Computing systems are merging**
 - Embedded systems were once stand-alone
 - Now modern limousines are distributed systems on wheels
 - Workstations were once networked systems
 - They now use parallel hardware (e.g. SMP, GPU)
 - Transparency is increasing (e.g. peer-to-peer)
 - Distributed systems were once local
 - Now the web is the computer (SETI@Home)
 - Parallel systems were once run on supercomputers
 - Clusters are now made of off-the-shelf computers with high-speed buses and networks
- **Operating systems are being challenged**
 - Light enough to support a stand-alone system
 - Powerful enough to support a distributed system
 - And parallelism on both cases

Distributed Systems

- Set of loosely coupled computers interconnected by a network
- Each computer has its own **local** resources plus **remote** resources from other computers in the set
- Processes on a distributed system access resources independently of whether they are local or remote (**location transparency**)
- Inter-process communication is mostly based on **message passing**
- Process models
 - Client-Server
 - Server has a resource that is used by the client
 - Peer-to-Peer
 - Both partner processes share some of their resources

Motivation

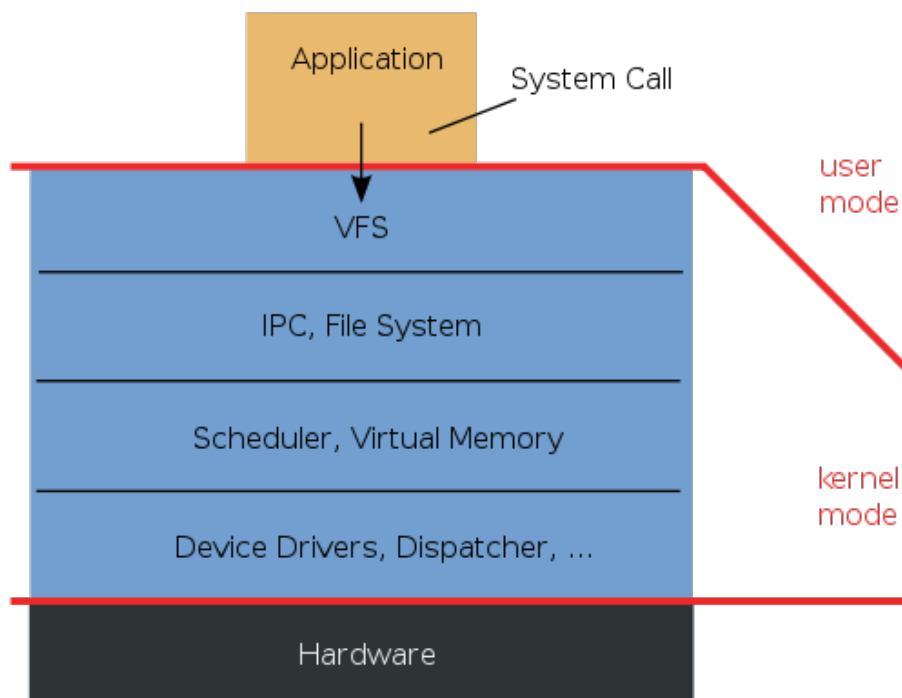
- Resource **sharing**
 - Remote file sharing, printing, access to special devices (scanner, CD writer, etc)
 - Distributed databases
- Computation **speedup**
 - Tasks can be partitioned and distributed
- **Reliability**
 - The failure of a node does not necessarily disrupts the system
- **Scalability**
 - New nodes can be aggregated to the system on demand
- Pitfalls: **complexity and security**

Transparency

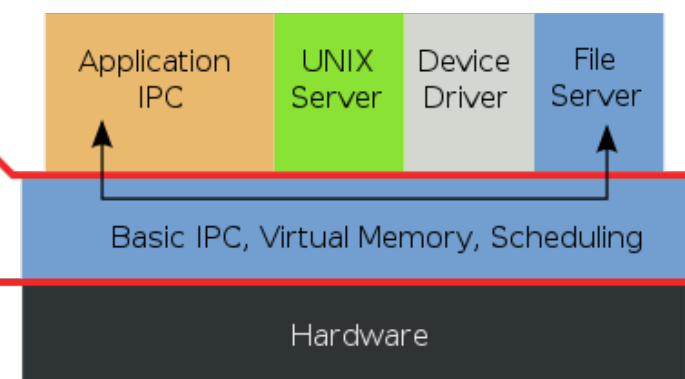
- **Location** transparency
 - Local and remote objects look just the same
 - No need to specify location
- **Migration** transparency
 - Objects change location, their names are preserved
- **Replication** transparency
 - Objects can be automatically replicated (consistency)
- **Concurrency** transparency
 - Objects can be concurrently manipulated without explicit synchronization
- **Parallelism** transparency
 - Automatic parallelization

μ -kernels and Distributed Systems

Monolithic Kernel based Operating System

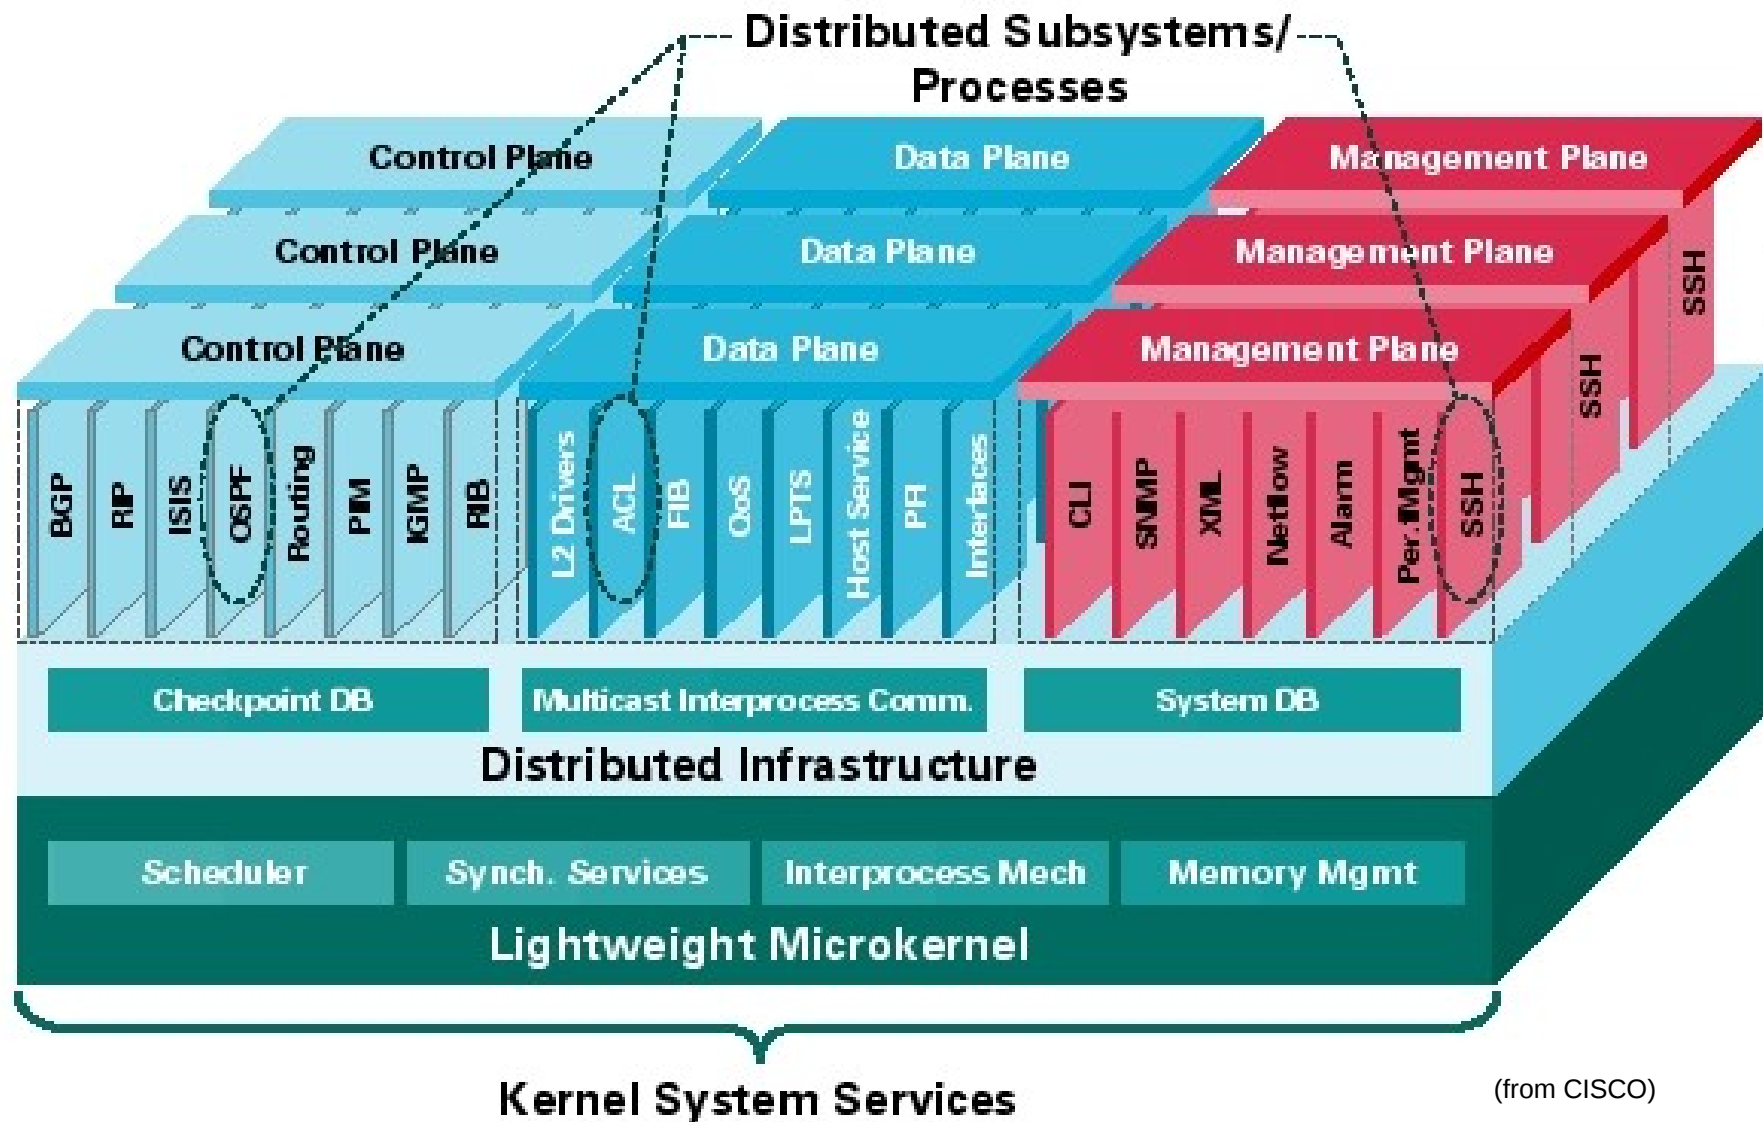


Microkernel based Operating System




(from Wikipedia)

CISCO IOS XR



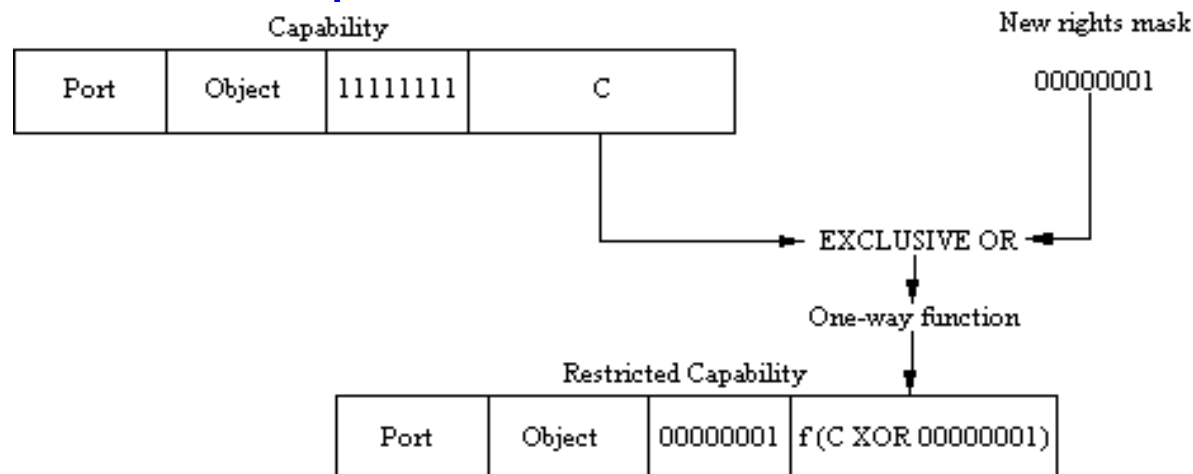
(from CISCO)

Remote Procedure Call (RPC)

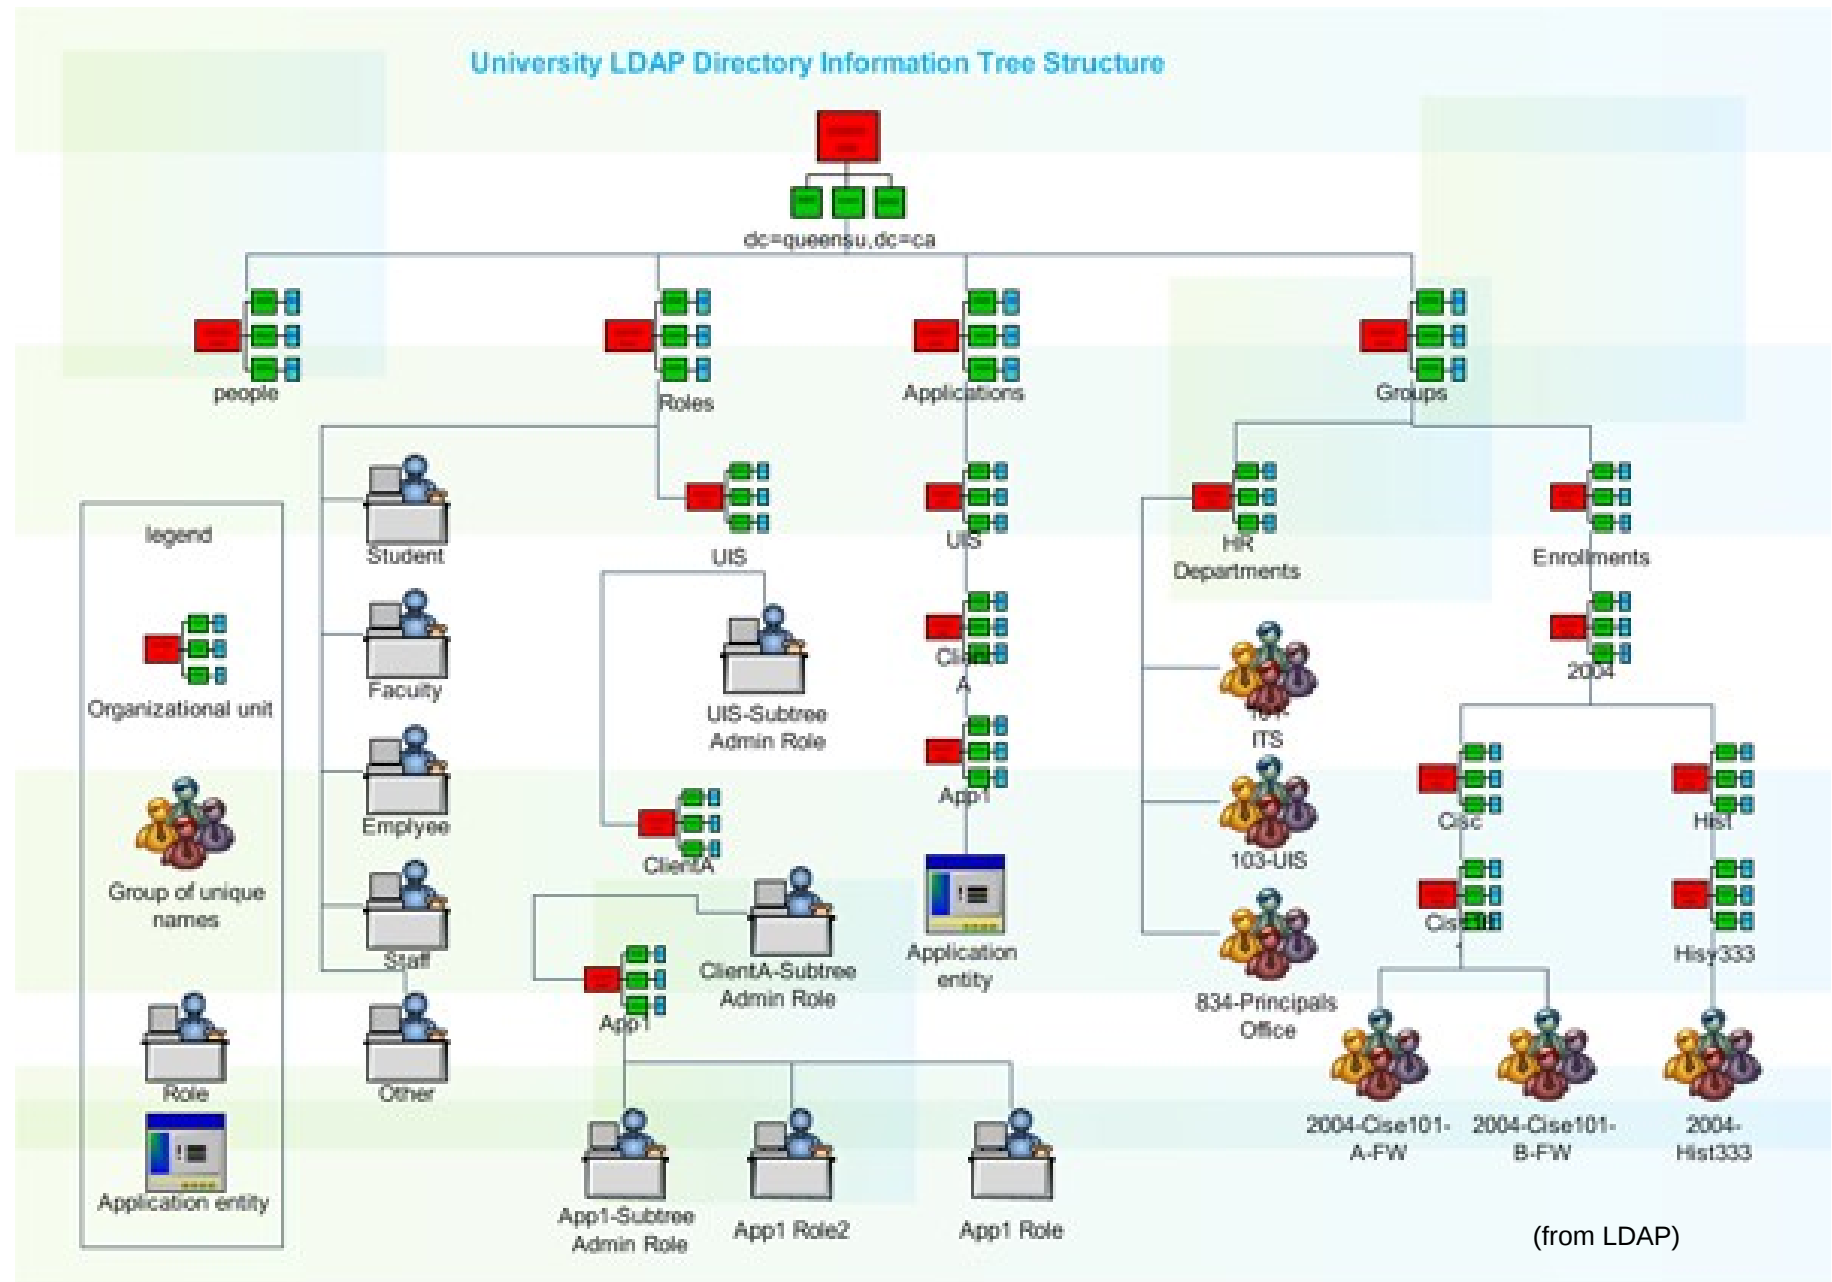


Object Identification

- Implicit id (domain specific)
 - Object **pointer**
- Local id
 - Object **counter** with reuse and overflow control
- Global id
 - **Host** id + local id
 - Easy on a MAC-assigned, IP-based world
- Capability
 - Global id + **permissions** + **secret**




Name Server




Communication Patterns


unicast




broadcast



multicast

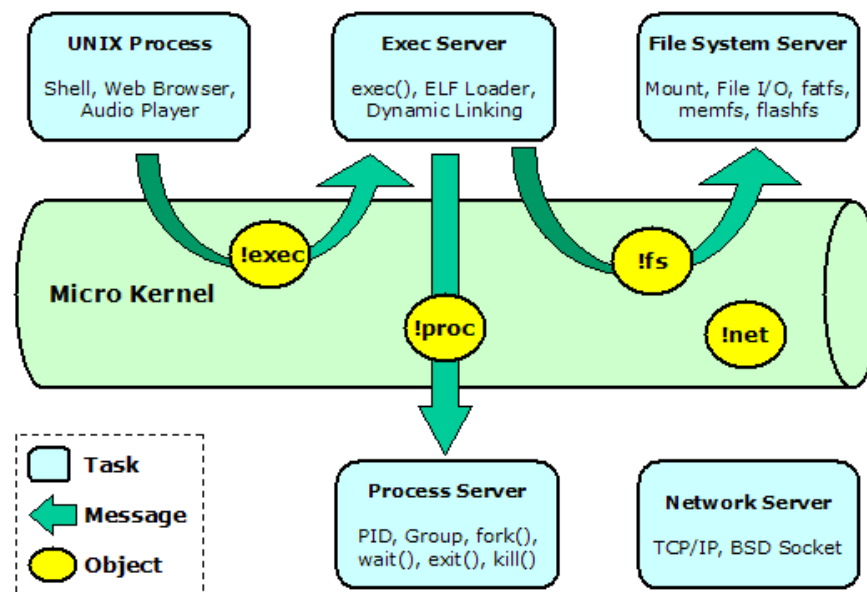


anycast



(from Wikipedia)

Inter-process Communication



- Messages used to request u-kernel services can be **forwarded** to other hosts
 - Global id
 - Network driver/service
- **Foundation** for all other distributed system services